

# ECE565 Homework1 Report

Yisong Zou  
NetId: yz558

## 1. (30 points) Code Optimization – “Beat the Compiler”.

### (a)

Results when using gcc -O2 -o loop\_performance loop\_performance.c

#### gcc -O2

Ten Runs	Argument (10000000) (seven 0s)	Argument (100000000) (eight 0s)
1st(millisecons)	27.322	268.833
2nd(millisecons)	26.391	258.334
3rd(millisecons)	26.624	258.541
4th(millisecons)	26.996	268.294
5th(millisecons)	27.273	259.977
6th(millisecons)	28.181	266.166
7th(millisecons)	26.739	264.944
8th(millisecons)	26.963	266.627
9th(millisecons)	26.789	264.117
10th(millisecons)	26.486	267.082
Shortest Time	26.391	258.334

Results when using gcc -O3 -o loop\_performance loop\_performance.c

#### gcc -O3

Ten Runs	Argument (10000000) (seven 0s)	Argument (100000000) (eight 0s)
1st(millisecons)	21.923	214.149
2nd(millisecons)	23.527	208.560
3rd(millisecons)	22.039	216.123
4th(millisecons)	22.158	207.204
5th(millisecons)	21.986	217.395
6th(millisecons)	21.906	212.907
7th(millisecons)	21.450	216.722
8th(millisecons)	21.212	212.022
9th(millisecons)	21.891	207.504
10th(millisecons)	21.978	215.260
Shortest Time	21.212	207.204

**(b)** Describe the type of machine and environment you are running on

- 1) processor architecture: x86-64
- 2) CPU frequency: CPU MHz: 2300.000
- 3) OS type: Ubuntu 18.04.5 LTS
- 4) It is inside duke VM

**(c)** Study the code in the do\_loops() function.

## Loop Unrolling

Code:

```
1. void do_loops(int *a, int *b, int *c, int N) {
2.     int i;
3.     for (i = N - 1; i >= 1; i -= 9) {
4.         a[i] = a[i] + 1;
5.         a[i - 1] = a[i - 1] + 1;
6.         a[i - 2] = a[i - 2] + 1;
7.         a[i - 3] = a[i - 3] + 1;
8.         a[i - 4] = a[i - 4] + 1;
9.         a[i - 5] = a[i - 5] + 1;
10.        a[i - 6] = a[i - 6] + 1;
11.        a[i - 7] = a[i - 7] + 1;
12.        a[i - 8] = a[i - 8] + 1;
13.    }
14.    for (i = 1; i < N; i += 9) {
15.        b[i] = a[i + 1] + 3;
16.        b[i + 1] = a[i + 2] + 3;
17.        b[i + 2] = a[i + 3] + 3;
18.        b[i + 3] = a[i + 4] + 3;
19.        b[i + 4] = a[i + 5] + 3;
20.        b[i + 5] = a[i + 6] + 3;
21.        b[i + 6] = a[i + 7] + 3;
22.        b[i + 7] = a[i + 8] + 3;
23.        b[i + 8] = a[i + 9] + 3;
24.    }
25.    for (i = 1; i < N; i += 9) {
26.        c[i] = b[i - 1] + 2;
27.        c[i + 1] = b[i] + 2;
28.        c[i + 2] = b[i + 1] + 2;
29.        c[i + 3] = b[i + 2] + 2;
30.        c[i + 4] = b[i + 3] + 2;
```

```

31.     c[i + 5] = b[i + 4] + 2;
32.     c[i + 6] = b[i + 5] + 2;
33.     c[i + 7] = b[i + 6] + 2;
34.     c[i + 8] = b[i + 7] + 2;
35. }
36. }

```

### gcc -O2 with Loop Unrolling

Ten Runs	Argument (10000000) (seven 0s)	Argument (100000000) (eight 0s)
1st(millisecons)	23.937	235.889
2nd(millisecons)	24.946	238.672
3rd(millisecons)	23.590	226.480
4th(millisecons)	24.058	222.157
5th(millisecons)	23.796	236.752
6th(millisecons)	23.746	227.602
7th(millisecons)	24.344	225.330
8th(millisecons)	24.459	226.794
9th(millisecons)	24.585	232.023
10th(millisecons)	24.009	227.108
Shortest Time	23.590	222.157

### gcc -O3 with Loop Unrolling

Ten Runs	Argument (10000000) (seven 0s)	Argument (100000000) (eight 0s)
1st(millisecons)	24.180	226.786
2nd(millisecons)	23.978	227.899
3rd(millisecons)	24.409	226.480
4th(millisecons)	24.450	226.634
5th(millisecons)	23.796	227.372
6th(millisecons)	23.746	227.602
7th(millisecons)	24.176	227.249
8th(millisecons)	23.175	227.116
9th(millisecons)	24.262	226.546
10th(millisecons)	23.702	227.108
Shortest Time	23.175	226.480

**Discuss:** (Does it match your expectation? What reason(s) might explain the result?)

According to the output of objdump as the follows, the unrolled version has more instructions done in a single loop and this will result in fewer loop management instructions. This matches my expectation, as unrolling reduce instruction count by make fewer loops, then there will be less loop management

instructions, it will run faster in gcc -O2. However, as -O3 has already have this optimization, it does not show improvement, and seems make it a little worse.

objdump results for original code(-O2)(Each color block represents a loop)

1.	a75:	7e 2d	jle	aa4 <do_loops+0x34>
2.	a77:	48 98	cltq	
3.	a79:	44 8d 49 fe	lea	-0x2(%rcx),%r9d
4.	a7d:	4c 8d 04 85 00 00 00	lea	0x0(,%rax,4),%r8
5.	a84:	00		
6.	a85:	49 c1 e1 02	shl	\$0x2,%r9
7.	a89:	4a 8d 04 07	lea	(%rdi,%r8,1),%rax
8.	a8d:	4e 8d 44 07 fc	lea	-0x4(%rdi,%r8,1),%r8
9.	a92:	4d 29 c8	sub	%r9,%r8
10.	a95:	0f 1f 00	nopl	(%rax)
11.	a98:	83 00 01	addl	\$0x1,(%rax)
12.	a9b:	48 83 e8 04	sub	\$0x4,%rax
13.	a9f:	4c 39 c0	cmp	%r8,%rax
14.	aa2:	75 f4	jne	a98 <do_loops+0x28>
15.	aa4:	83 f9 01	cmp	\$0x1,%ecx
16.	aa7:	7e 4a	jle	af3 <do_loops+0x83>
17.	aa9:	8d 41 fe	lea	-0x2(%rcx),%eax
18.	aac:	4c 8d 04 85 04 00 00	lea	0x4(,%rax,4),%r8
19.	ab3:	00		
20.	ab4:	31 c0	xor	%eax,%eax
21.	ab6:	66 2e 0f 1f 84 00 00	nopw	%cs:0x0(%rax,%rax,1)
22.	abd:	00 00 00		
23.	ac0:	8b 4c 07 08	mov	0x8(%rdi,%rax,1),%ecx
24.	ac4:	83 c1 03	add	\$0x3,%ecx
25.	ac7:	89 4c 06 04	mov	%ecx,0x4(%rsi,%rax,1)
26.	acb:	48 83 c0 04	add	\$0x4,%rax
27.	acf:	49 39 c0	cmp	%rax,%r8
28.	ad2:	75 ec	jne	ac0 <do_loops+0x50>
29.	ad4:	31 c9	xor	%ecx,%ecx
30.	ad6:	66 2e 0f 1f 84 00 00	nopw	%cs:0x0(%rax,%rax,1)
31.	add:	00 00 00		
32.	ae0:	8b 3c 0e	mov	(%rsi,%rcx,1),%edi
33.	ae3:	83 c7 02	add	\$0x2,%edi
34.	ae6:	89 7c 0a 04	mov	%edi,0x4(%rdx,%rcx,1)
35.	aea:	48 83 c1 04	add	\$0x4,%rcx
36.	aee:	48 39 c1	cmp	%rax,%rcx
37.	af1:	75 ed	jne	ae0 <do_loops+0x70>

## objdump results for unrolling code(-O2) (Each color block represents a loop)

```
1.  a88:  41 83 e8 09          sub     $0x9,%r8d
2.  a8c:  83 40 20 01          addl    $0x1,0x20(%rax)
3.  a90:  83 40 1c 01          addl    $0x1,0x1c(%rax)
4.  a94:  83 40 18 01          addl    $0x1,0x18(%rax)
5.  a98:  83 40 14 01          addl    $0x1,0x14(%rax)
6.  a9c:  83 40 10 01          addl    $0x1,0x10(%rax)
7.  aa0:  83 40 0c 01          addl    $0x1,0xc(%rax)
8.  aa4:  83 40 08 01          addl    $0x1,0x8(%rax)
9.  aa8:  83 40 04 01          addl    $0x1,0x4(%rax)
10. aac:  83 00 01             addl    $0x1,(%rax)
11. aaf:  48 83 e8 24          sub     $0x24,%rax
12. ab3:  45 85 c0             test    %r8d,%r8d
13. ab6:  7f d0               jg      a88 <do_loops+0x18>
14. ab8:  83 f9 01             cmp     $0x1,%ecx
15. abb:  0f 8e f6 00 00 00    jle     bb7 <do_loops+0x147>
16. ac1:  48 8d 46 04          lea     0x4(%rsi),%rax
17. ac5:  48 83 c7 08          add     $0x8,%rdi
18. ac9:  41 b9 01 00 00 00    mov     $0x1,%r9d
19. acf:  90                  nop
20. ad0:  44 8b 17             mov     (%rdi),%r10d
21. ad3:  41 83 c1 09          add     $0x9,%r9d
22. ad7:  48 83 c7 24          add     $0x24,%rdi
23. adb:  48 83 c0 24          add     $0x24,%rax
24. adf:  45 8d 42 03          lea     0x3(%r10),%r8d
25. ae3:  44 89 40 dc          mov     %r8d,-0x24(%rax)
26. ae7:  44 8b 5f e0          mov     -0x20(%rdi),%r11d
27. aeb:  45 8d 43 03          lea     0x3(%r11),%r8d
28. aef:  44 89 40 e0          mov     %r8d,-0x20(%rax)
29. af3:  44 8b 57 e4          mov     -0x1c(%rdi),%r10d
30. af7:  45 8d 42 03          lea     0x3(%r10),%r8d
31. afb:  44 89 40 e4          mov     %r8d,-0x1c(%rax)
32. aff:  44 8b 5f e8          mov     -0x18(%rdi),%r11d
33. b03:  45 8d 43 03          lea     0x3(%r11),%r8d
34. b07:  44 89 40 e8          mov     %r8d,-0x18(%rax)
35. b0b:  44 8b 57 ec          mov     -0x14(%rdi),%r10d
36. b0f:  45 8d 42 03          lea     0x3(%r10),%r8d
37. b13:  44 89 40 ec          mov     %r8d,-0x14(%rax)
38. b17:  44 8b 5f f0          mov     -0x10(%rdi),%r11d
39. b1b:  45 8d 43 03          lea     0x3(%r11),%r8d
40. b1f:  44 89 40 f0          mov     %r8d,-0x10(%rax)
41. b23:  44 8b 57 f4          mov     -0xc(%rdi),%r10d
42. b27:  45 8d 42 03          lea     0x3(%r10),%r8d
43. b2b:  44 89 40 f4          mov     %r8d,-0xc(%rax)
```

44.	b2f:	44 8b 5f f8	mov	-0x8(%rdi),%r11d
45.	b33:	45 8d 43 03	lea	0x3(%r11),%r8d
46.	b37:	44 89 40 f8	mov	%r8d,-0x8(%rax)
47.	b3b:	44 8b 57 fc	mov	-0x4(%rdi),%r10d
48.	b3f:	45 8d 42 03	lea	0x3(%r10),%r8d
49.	b43:	44 89 40 fc	mov	%r8d,-0x4(%rax)
50.	b47:	44 39 c9	cmp	%r9d,%ecx
51.	b4a:	7f 84	jg	ad0 <do_loops+0x60>
52.	b4c:	48 83 c2 04	add	\$0x4,%rdx
53.	b50:	bf 01 00 00 00	mov	\$0x1,%edi
54.	b55:	0f 1f 00	nopl	(%rax)
55.	b58:	8b 06	mov	(%rsi),%eax
56.	b5a:	83 c7 09	add	\$0x9,%edi
57.	b5d:	48 83 c6 24	add	\$0x24,%rsi
58.	b61:	48 83 c2 24	add	\$0x24,%rdx
59.	b65:	83 c0 02	add	\$0x2,%eax
60.	b68:	89 42 dc	mov	%eax,-0x24(%rdx)
61.	b6b:	8b 46 e0	mov	-0x20(%rsi),%eax
62.	b6e:	83 c0 02	add	\$0x2,%eax
63.	b71:	89 42 e0	mov	%eax,-0x20(%rdx)
64.	b74:	8b 46 e4	mov	-0x1c(%rsi),%eax
65.	b77:	83 c0 02	add	\$0x2,%eax
66.	b7a:	89 42 e4	mov	%eax,-0x1c(%rdx)
67.	b7d:	8b 46 e8	mov	-0x18(%rsi),%eax
68.	b80:	83 c0 02	add	\$0x2,%eax
69.	b83:	89 42 e8	mov	%eax,-0x18(%rdx)
70.	b86:	8b 46 ec	mov	-0x14(%rsi),%eax
71.	b89:	83 c0 02	add	\$0x2,%eax
72.	b8c:	89 42 ec	mov	%eax,-0x14(%rdx)
73.	b8f:	8b 46 f0	mov	-0x10(%rsi),%eax
74.	b92:	83 c0 02	add	\$0x2,%eax
75.	b95:	89 42 f0	mov	%eax,-0x10(%rdx)
76.	b98:	8b 46 f4	mov	-0xc(%rsi),%eax
77.	b9b:	83 c0 02	add	\$0x2,%eax
78.	b9e:	89 42 f4	mov	%eax,-0xc(%rdx)
79.	ba1:	8b 46 f8	mov	-0x8(%rsi),%eax
80.	ba4:	83 c0 02	add	\$0x2,%eax
81.	ba7:	89 42 f8	mov	%eax,-0x8(%rdx)
82.	baa:	8b 46 fc	mov	-0x4(%rsi),%eax
83.	bad:	83 c0 02	add	\$0x2,%eax
84.	bb0:	89 42 fc	mov	%eax,-0x4(%rdx)
85.	bb3:	39 f9	cmp	%edi,%ecx
86.	bb5:	7f a1	jg	b58 <do_loops+0xe8>

# Loop Fusion

Code:

```
1. void do_loops(int *a, int *b, int *c, int N) {  
2.     int i;  
3.     for (i = N - 1; i >= 1; i--) {  
4.         a[i] = a[i] + 1;  
5.     }  
6.     for (i = 1; i < N; i++) {  
7.         b[i] = a[i + 1] + 3;  
8.         c[i] = b[i - 1] + 2;  
9.     }  
10. }
```

## gcc -O2 with Loop Fusion

Argument(10000000) (seven 0s) (Shortest 25.574)

```
1. Time=25.574000 milliseconds  
2. Time=26.018000 milliseconds  
3. Time=25.936000 milliseconds  
4. Time=26.117000 milliseconds  
5. Time=26.006000 milliseconds  
6. Time=25.869000 milliseconds  
7. Time=25.853000 milliseconds  
8. Time=26.084000 milliseconds  
9. Time=25.948000 milliseconds  
10. Time=25.891000 milliseconds
```

## gcc -O2 with Loop Fusion

Argument(100000000) (eight 0s) (Shortest 263.761)

```
1. Time=272.514000 milliseconds  
2. Time=266.591000 milliseconds  
3. Time=263.761000 milliseconds  
4. Time=265.740000 milliseconds  
5. Time=268.265000 milliseconds  
6. Time=268.219000 milliseconds  
7. Time=265.682000 milliseconds  
8. Time=275.816000 milliseconds  
9. Time=264.860000 milliseconds  
10. Time=274.635000 milliseconds
```

### gcc -O3 with Loop Fusion

Argument (10000000) (seven 0s) (Shortest 31.054)

```
1. Time=31.582000 milliseconds
2. Time=32.287000 milliseconds
3. Time=31.580000 milliseconds
4. Time=31.054000 milliseconds
5. Time=31.968000 milliseconds
6. Time=31.666000 milliseconds
7. Time=32.605000 milliseconds
8. Time=31.840000 milliseconds
9. Time=31.577000 milliseconds
10. Time=32.022000 milliseconds
```

### gcc -O3 with Loop Fusion

Argument (100000000) (eight 0s) (Shortest 306.348)

```
1. Time=331.303000 milliseconds
2. Time=314.896000 milliseconds
3. Time=306.348000 milliseconds
4. Time=313.078000 milliseconds
5. Time=331.573000 milliseconds
6. Time=312.558000 milliseconds
7. Time=323.960000 milliseconds
8. Time=317.125000 milliseconds
9. Time=310.286000 milliseconds
10. Time=314.238000 milliseconds
```

**Discuss:** (Does it match your expectation? What reason(s) might explain the result?)

According to the objdump results as the following, this matches my expectation, as fusion reduce instruction count by make fewer loop management instructions, it should run faster. For -O2 it is faster and for -O3 it is much slower, I think this might because -O3 has better optimization than fusion or it already has fusion.

objdump results for fusion code(-O2) (Each color block represents a loop)

```
1. 000000000000a70 <do_loops>:
2. a70: 8d 41 ff          lea    -0x1(%rcx),%eax
3. a73: 85 c0            test   %eax,%eax
4. a75: 7e 2d            jle    aa4 <do_loops+0x34>
5. a77: 48 98            cltq
6. a79: 44 8d 49 fe      lea    -0x2(%rcx),%r9d
7. a7d: 4c 8d 04 85 00 00 00 lea    0x0(,%rax,4),%r8
```



8.	a84:	00		
9.	a85:	49 c1 e1 02	shl	\$0x2,%r9
10.	a89:	4a 8d 04 07	lea	(%rdi,%r8,1),%rax
11.	a8d:	4e 8d 44 07 fc	lea	-0x4(%rdi,%r8,1),%r8
12.	a92:	4d 29 c8	sub	%r9,%r8
13.	a95:	0f 1f 00	nopl	(%rax)
14.	a98:	83 00 01	addl	\$0x1,(%rax)
15.	a9b:	48 83 e8 04	sub	\$0x4,%rax
16.	a9f:	4c 39 c0	cmp	%r8,%rax
17.	aa2:	75 f4	jne	a98 <do_loops+0x28>
18.	aa4:	83 f9 01	cmp	\$0x1,%ecx
19.	aa7:	7e 34	jle	add <do_loops+0x6d>
20.	aa9:	8d 41 fe	lea	-0x2(%rcx),%eax
21.	aac:	4c 8d 04 85 08 00 00	lea	0x8(,%rax,4),%r8
22.	ab3:	00		
23.	ab4:	b8 04 00 00 00	mov	\$0x4,%eax
24.	ab9:	0f 1f 80 00 00 00 00	nopl	0x0(%rax)
25.	ac0:	8b 4c 07 04	mov	0x4(%rdi,%rax,1),%ecx
26.	ac4:	83 c1 03	add	\$0x3,%ecx
27.	ac7:	89 0c 06	mov	%ecx,(%rsi,%rax,1)
28.	aca:	8b 4c 06 fc	mov	-0x4(%rsi,%rax,1),%ecx
29.	ace:	83 c1 02	add	\$0x2,%ecx
30.	ad1:	89 0c 02	mov	%ecx,(%rdx,%rax,1)
31.	ad4:	48 83 c0 04	add	\$0x4,%rax
32.	ad8:	4c 39 c0	cmp	%r8,%rax
33.	adb:	75 e3	jne	ac0 <do_loops+0x50>
34.	add:	f3 c3	repz retq	
35.	adf:	90	nop	

## Loop Strip Mining

Code:

```

1. void do_loops(int *a, int *b, int *c, int N) {
2.     int i;
3.     int j;
4.     for (j = N - 1; j >= 1; j -= 9) {
5.         for (i = j; i > j - 9; i--) {
6.             a[i] = a[i] + 1;
7.         }
8.     }
9.     for (j = 1; j < N; j += 9) {
10.        for (i = j; i < j + 9; i++) {

```

```

11.      b[i] = a[i + 1] + 3;
12.  }
13. }
14.  for (j = 1; j < N; j += 9) {
15.      for (i = j; i < j + 9; i++) {
16.          c[i] = b[i - 1] + 2;
17.      }
18.  }
19. }

```

### gcc -O2 with Loop Strip Mining

Argument(10000000) (seven 0s) (Shortest 31.956)

```

1. Time=32.337000 milliseconds
2. Time=32.175000 milliseconds
3. Time=32.877000 milliseconds
4. Time=32.511000 milliseconds
5. Time=32.330000 milliseconds
6. Time=32.358000 milliseconds
7. Time=32.070000 milliseconds
8. Time=32.057000 milliseconds
9. Time=32.152000 milliseconds
10. Time=31.956000 milliseconds

```

### gcc -O2 with Loop Strip Mining

Argument(100000000) (eight 0s) (Shortest 313.814)

```

1. Time=319.009000 milliseconds
2. Time=319.301000 milliseconds
3. Time=329.217000 milliseconds
4. Time=315.710000 milliseconds
5. Time=313.814000 milliseconds
6. Time=321.578000 milliseconds
7. Time=317.843000 milliseconds
8. Time=316.670000 milliseconds
9. Time=315.294000 milliseconds
10. Time=314.180000 milliseconds

```

### gcc -O3 with Loop Strip Mining

Argument(10000000) (seven 0s) (Shortest 22.863)

```

1. Time=23.165000 milliseconds

```

```

2. Time=23.554000 milliseconds
3. Time=23.611000 milliseconds
4. Time=23.617000 milliseconds
5. Time=23.122000 milliseconds
6. Time=24.354000 milliseconds
7. Time=22.863000 milliseconds
8. Time=27.051000 milliseconds
9. Time=23.004000 milliseconds
10. Time=23.087000 milliseconds

```

### **gcc -O3 with Loop Strip Mining**

Argument (100000000) (eight 0s) (Shortest 222.123)

```

1. Time=230.581000 milliseconds
2. Time=235.131000 milliseconds
3. Time=233.066000 milliseconds
4. Time=233.635000 milliseconds
5. Time=223.609000 milliseconds
6. Time=222.737000 milliseconds
7. Time=222.123000 milliseconds
8. Time=225.214000 milliseconds
9. Time=222.586000 milliseconds
10. Time=223.465000 milliseconds

```

**Discuss:** (Does it match your expectation? What reason(s) might explain the result?)

According to the objdump result(too long so I will just describe it here), this matches my expectation, as under this circumstance, mining cannot benefit from vectorization, it only adds many useless loop management instructions and more level of useless loops. For -O3 it is a little slower and for -O2 it is much slower, this is because this mining lead to more overhead of loop management instructions.

**(d)** Discuss whether you can beat the compiler

From the (C) part, I can beat the computer only in -O2 with loop unrolling, the statistics are as the following.

## ***My Code Results with Unrolling***

### **gcc -O2 with Loop Unrolling**

Ten Runs	Argument (100000000) (seven 0s)	Argument (100000000) (eight 0s)
1st(milliseconds)	23.937	235.889
2nd(milliseconds)	24.946	238.672

3rd(millisecond)	23.590	226.480
4th(millisecond)	24.058	222.157
5th(millisecond)	23.796	236.752
6th(millisecond)	23.746	227.602
7th(millisecond)	24.344	225.330
8th(millisecond)	24.459	226.794
9th(millisecond)	24.585	232.023
10th(millisecond)	24.009	227.108
Shortest Time	23.590	222.157

## ***Original Code Results***


### **gcc -O2**

Ten Runs	Argument (10000000) (seven 0s)	Argument (100000000) (eight 0s)
1st(millisecond)	27.322	268.833
2nd(millisecond)	26.391	258.334
3rd(millisecond)	26.624	258.541
4th(millisecond)	26.996	268.294
5th(millisecond)	27.273	259.977
6th(millisecond)	28.181	266.166
7th(millisecond)	26.739	264.944
8th(millisecond)	26.963	266.627
9th(millisecond)	26.789	264.117
10th(millisecond)	26.486	267.082
Shortest Time	26.391	258.334

**From the above statistics, my code have beaten the computer in -O2 option.**

## 2. (20 points) Dependence Analysis.

(a) Draw ITG



(b) List all the dependencies

(b). Dependences

$S_1[i, j] \Rightarrow T S_2[i+1, j+1]$  (loop carried)

$S_1[i, j] \Rightarrow T S_3[i, j]$  (loop independent)

$S_1[i-1, j+1] \Rightarrow A S_1[i, j]$  (loop carried)


$S_1[i, j] \Rightarrow A S_2[i, j]$  (loop independent)

$S_3[i-1, j] \Rightarrow T S_1[i, j]$  (loop carried)

$S_3[i-1, j] \Rightarrow T S_2[i+1, j]$  (loop carried)

$S_4[i, j] \Rightarrow T S_4[i+1, j-1]$  (loop carried)

(c) Draw the LDG



### 3. (20 points) Function In-lining and performance

(a).

**Without Inlining Function Code:**

```
__attribute__((noinline)) int add(int a, int b) { return (a + b); }
```

**Without inlining Original code:**

Argument(10000000) (seven 0s) (Shortest 28.311)

```
1. Time=28.357 milliseconds
2. Time=29.208 milliseconds
3. Time=28.336 milliseconds
4. Time=28.42 milliseconds
5. Time=28.405 milliseconds
6. Time=28.311 milliseconds
7. Time=28.441 milliseconds
8. Time=28.571 milliseconds
9. Time=28.46 milliseconds
10. Time=28.375 milliseconds
```

**Without inlining Original code:**

Argument(100000000) (eight 0s) (Shortest 281.478)

```
1. Time=282.653 milliseconds
2. Time=283.631 milliseconds
3. Time=282.651 milliseconds
4. Time=281.774 milliseconds
5. Time=282.116 milliseconds
6. Time=282.202 milliseconds
7. Time=282.153 milliseconds
8. Time=284.597 milliseconds
9. Time=281.881 milliseconds
10. Time=281.478 milliseconds
```

**Inlining Function Code:**

```
inline __attribute__((always_inline)) int add(int a, int b) { return (a + b)
; }
```

**With inlining:**

Argument(10000000) (seven 0s) (Shortest 10.169)

```

1. Time=10.193 milliseconds
2. Time=10.169 milliseconds
3. Time=10.281 milliseconds
4. Time=10.281 milliseconds
5. Time=10.609 milliseconds
6. Time=10.291 milliseconds
7. Time=10.782 milliseconds
8. Time=10.484 milliseconds
9. Time=10.268 milliseconds
10. Time=10.314 milliseconds

```

### With inlining:

Argument (100000000) (eight 0s) (Shortest 103.273)

```

1. Time=104.732 milliseconds
2. Time=104.525 milliseconds
3. Time=104.213 milliseconds
4. Time=105.15 milliseconds
5. Time=104.727 milliseconds
6. Time=103.273 milliseconds
7. Time=104.606 milliseconds
8. Time=104.956 milliseconds
9. Time=104.669 milliseconds
10. Time=104.098 milliseconds

```

(b)

We can see from the following two different assemblies that inline will cause the main function to directly contain the lines of the function while no\_inline version will make a call to the function.

### Assembly without inlining

1. e12:	e8 19 fd ff ff	callq	b30 <gettimeofday@plt>
2. e17:	44 89 e0	mov	%r12,%eax
3. e1a:	4c 8b 54 24 50	mov	0x50(%rsp),%r10
4. e1f:	4c 8b 4c 24 30	mov	0x30(%rsp),%r9
5. e24:	4c 8b 44 24 70	mov	0x70(%rsp),%r8
6. e29:	48 8d 0c 85 04 00 00	lea	0x4(,%rax,4),%rcx
7. e30:	00		
8. e31:	31 d2	xoqr	%edx,%edx
9. e33:	0f 1f 44 00 00	nopl	0x0(%rax,%rax,1)
10. e38:	41 8b 34 12	mov	(%r10,%rdx,1),%esi
11. e3c:	41 8b 3c 11	mov	(%r9,%rdx,1),%edi
12. e40:	e8 eb 04 00 00	callq	1330 <_Z3addii>



13. e45:	41 89 04 10	mov	%eax, (%r8,%rdx,1)
14. e49:	48 83 c2 04	add	\$0x4,%rdx
15. e4d:	48 39 ca	cmp	%rcx,%rdx
16. e50:	75 e6	jne	e38 <main+0x2c8>
17. e52:	31 f6	xor	%esi,%esi
18. e54:	48 89 ef	mov	%rbp,%rdi
19. e57:	e8 d4 fc ff ff	callq	b30 <gettimeofday@plt>

## Assembly with inlining

1. e12:	e8 19 fd ff ff	callq	b30 <gettimeofday@plt>
2. e17:	48 8b 7c 24 50	mov	0x50(%rsp),%rdi
3. e1c:	48 8b 74 24 70	mov	0x70(%rsp),%rsi
4. e21:	4c 8b 44 24 30	mov	0x30(%rsp),%r8
5. e26:	48 8d 47 10	lea	0x10(%rdi),%rax
6. e2a:	48 8d 4e 10	lea	0x10(%rsi),%rcx
7. e2e:	48 39 c6	cmp	%rax,%rsi
8. e31:	0f 93 c2	setae	%dl
9. e34:	48 39 cf	cmp	%rcx,%rdi
10. e37:	0f 93 c0	setae	%al
11. e3a:	09 c2	or	%eax,%edx
12. e3c:	49 8d 40 10	lea	0x10(%r8),%rax
13. e40:	48 39 c6	cmp	%rax,%rsi
14. e43:	0f 93 c0	setae	%al
15. e46:	49 39 c8	cmp	%rcx,%r8
16. e49:	0f 93 c1	setae	%cl
17. e4c:	09 c8	or	%ecx,%eax
18. e4e:	84 c2	test	%al,%dl
19. e50:	0f 84 7e 03 00 00	je	11d4 <main+0x664>
20. e56:	83 fb 08	cmp	\$0x8,%ebx
21. e59:	0f 86 75 03 00 00	jbe	11d4 <main+0x664>
22. e5f:	48 89 f9	mov	%rdi,%rcx
23. e62:	48 c1 e9 02	shr	\$0x2,%rcx
24. e66:	48 f7 d9	neg	%rcx
25. e69:	83 e1 03	and	\$0x3,%ecx
26. e6c:	0f 84 5a 03 00 00	je	11cc <main+0x65c>
27. e72:	41 8b 00	mov	(%r8),%eax
28. e75:	03 07	add	(%rdi),%eax
29. e77:	83 f9 01	cmp	\$0x1,%ecx
30. e7a:	89 06	mov	%eax,(%rsi)
31. e7c:	0f 84 b0 03 00 00	je	1232 <main+0x6c2>
32. e82:	41 8b 40 04	mov	0x4(%r8),%eax
33. e86:	03 47 04	add	0x4(%rdi),%eax

34.	e89:	83 f9 02	cmp	\$0x2,%ecx
35.	e8c:	89 46 04	mov	%eax,0x4(%rsi)
36.	e8f:	0f 84 dc 03 00 00	je	1271 <main+0x701>
37.	e95:	41 8b 40 08	mov	0x8(%r8),%eax
38.	e99:	03 47 08	add	0x8(%rdi),%eax
39.	e9c:	41 b9 03 00 00 00	mov	\$0x3,%r9d
40.	ea2:	89 46 08	mov	%eax,0x8(%rsi)
41.	ea5:	41 89 db	mov	%ebx,%r11d
42.	ea8:	31 c0	xor	%eax,%eax
43.	eea:	31 d2	xor	%edx,%edx
44.	eac:	41 29 cb	sub	%ecx,%r11d
45.	eaf:	89 c9	mov	%ecx,%ecx
46.	eb1:	48 c1 e1 02	shl	\$0x2,%rcx
47.	eb5:	45 89 da	mov	%r11d,%r10d
48.	eb8:	4c 8d 34 0f	lea	(%rdi,%rcx,1),%r14
49.	ebc:	4d 8d 2c 08	lea	(%r8,%rcx,1),%r13
50.	ec0:	41 c1 ea 02	shr	\$0x2,%r10d
51.	ec4:	48 01 f1	add	%rsi,%rcx
52.	ec7:	66 0f 1f 84 00 00 00	nopw	0x0(%rax,%rax,1)
53.	ece:	00 00		
54.	ed0:	f3 41 0f 6f 44 05 00	movdqu	0x0(%r13,%rax,1),%xmm0
55.	ed7:	83 c2 01	add	\$0x1,%edx
56.	eda:	66 41 0f fe 04 06	padd	(%r14,%rax,1),%xmm0
57.	ee0:	0f 11 04 01	movups	%xmm0, (%rcx,%rax,1)
58.	ee4:	48 83 c0 10	add	\$0x10,%rax
59.	ee8:	41 39 d2	cmp	%edx,%r10d
60.	eeb:	77 e3	ja	ed0 <main+0x360>
61.	eed:	44 89 da	mov	%r11d,%edx
62.	ef0:	83 e2 fc	and	\$0xffffffffc,%edx
63.	ef3:	41 39 d3	cmp	%edx,%r11d
64.	ef6:	42 8d 04 0a	lea	(%rdx,%r9,1),%eax
65.	efa:	74 70	je	f6c <main+0x3fc>
66.	efc:	48 63 d0	movslq	%eax,%rdx
67.	eff:	41 8b 0c 90	mov	(%r8,%rdx,4),%ecx
68.	f03:	03 0c 97	add	(%rdi,%rdx,4),%ecx
69.	f06:	89 0c 96	mov	%ecx, (%rsi,%rdx,4)
70.	f09:	8d 50 01	lea	0x1(%rax),%edx
71.	f0c:	39 d3	cmp	%edx,%ebx
72.	f0e:	7e 5c	jle	f6c <main+0x3fc>
73.	f10:	48 63 d2	movslq	%edx,%rdx
74.	f13:	41 8b 0c 90	mov	(%r8,%rdx,4),%ecx
75.	f17:	03 0c 97	add	(%rdi,%rdx,4),%ecx
76.	f1a:	89 0c 96	mov	%ecx, (%rsi,%rdx,4)
77.	f1d:	8d 50 02	lea	0x2(%rax),%edx

78.	f20:	39 d3	cmp	%edx,%ebx
79.	f22:	7e 48	jle	f6c <main+0x3fc>
80.	f24:	48 63 d2	movslq	%edx,%rdx
81.	f27:	41 8b 0c 90	mov	(%r8,%rdx,4),%ecx
82.	f2b:	03 0c 97	add	(%rdi,%rdx,4),%ecx
83.	f2e:	89 0c 96	mov	%ecx,(%rsi,%rdx,4)
84.	f31:	8d 50 03	lea	0x3(%rax),%edx
85.	f34:	39 d3	cmp	%edx,%ebx
86.	f36:	7e 34	jle	f6c <main+0x3fc>
87.	f38:	48 63 d2	movslq	%edx,%rdx
88.	f3b:	41 8b 0c 90	mov	(%r8,%rdx,4),%ecx
89.	f3f:	03 0c 97	add	(%rdi,%rdx,4),%ecx
90.	f42:	89 0c 96	mov	%ecx,(%rsi,%rdx,4)
91.	f45:	8d 50 04	lea	0x4(%rax),%edx
92.	f48:	39 d3	cmp	%edx,%ebx
93.	f4a:	7e 20	jle	f6c <main+0x3fc>
94.	f4c:	48 63 d2	movslq	%edx,%rdx
95.	f4f:	83 c0 05	add	\$0x5,%eax
96.	f52:	41 8b 0c 90	mov	(%r8,%rdx,4),%ecx
97.	f56:	03 0c 97	add	(%rdi,%rdx,4),%ecx
98.	f59:	39 c3	cmp	%eax,%ebx
99.	f5b:	89 0c 96	mov	%ecx,(%rsi,%rdx,4)
100.	f5e:	7e 0c	jle	f6c <main+0x3fc>
101.	f60:	48 98	cltq	
102.	f62:	41 8b 14 80	mov	(%r8,%rax,4),%edx
103.	f66:	03 14 87	add	(%rdi,%rax,4),%edx
104.	f69:	89 14 86	mov	%edx,(%rsi,%rax,4)
105.	f6c:	31 f6	xor	%esi,%esi
106.	f6e:	48 89 ef	mov	%rbp,%rdi
107.	f71:	e8 ba fb ff ff	callq	b30 <gettimeofday@plt>

(c)

**The performance matches my expectation. There is a great improvement on the speed of the function because it avoids the function call & return instructions and will allow the compiler to better optimize the code.**

(d)

**The performance of the original code:**

Argument(10000000) (seven 0s) (Shortest 11.101)

1. Time=11.459 milliseconds
2. Time=11.826 milliseconds

```
3. Time=11.486 milliseconds
4. Time=12.155 milliseconds
5. Time=11.531 milliseconds
6. Time=11.983 milliseconds
7. Time=11.517 milliseconds
8. Time=12.101 milliseconds
9. Time=11.541 milliseconds
10. Time=12.048 milliseconds
```

Argument(100000000) (eight 0s) (Shortest 114.973)

```
1. Time=117.608 milliseconds
2. Time=115.398 milliseconds
3. Time=115.349 milliseconds
4. Time=114.973 milliseconds
5. Time=116.951 milliseconds
6. Time=117.818 milliseconds
7. Time=117.138 milliseconds
8. Time=117.996 milliseconds
9. Time=115.649 milliseconds
10. Time=116.269 milliseconds
```

**The original code's performance is nearly the same as the in-lining code, I think it is because -O3 compiler has already done the in-lining by default, so there is not such a large difference.**

## 4. (15 points) Loop Transformations

### Original Code:

```
1. ...
2. int a[N][4];
3. int rand_number = rand();
4. for (i=0; i<4; i++) {
5.     threshold = 2.0 * rand_number;
6.     for (j=0; j<N; j++) {
7.         if (threshold < 4) {
8.             sum = sum + a[j][i];
9.         } else {
10.            sum = sum + a[j][i] + 1;
11.        }
12.    }
13. } ...
```

### After Loop Invariant Hoisting

```
1. ...
2. int a[N][4];
3. int rand_number = rand();
4. threshold = 2.0 * rand_number;
5. for (i=0; i<4; i++) {
6.     for (j=0; j<N; j++) {
7.         if (threshold < 4) {
8.             sum = sum + a[j][i];
9.         } else {
10.            sum = sum + a[j][i] + 1;
11.        }
12.    }
13. } ...
```

### After Loop Unroll and Jam

```
1. ...
2. int a[N][4];
3. int rand_number = rand();
4. threshold = 2.0 * rand_number;
5. for (i=0; i<4; i+=4) {
6.     for (j=0; j<N; j++) {
7.         if (threshold < 4) {
8.             sum = sum + a[j][i];
```

```

9. sum = sum + a[j][i + 1];
10. sum = sum + a[j][i + 2];
11. sum = sum + a[j][i + 3];
12. } else {
13. sum = sum + a[j][i] + 1;
14. sum = sum + a[j][i + 1] + 1;
15. sum = sum + a[j][i + 2] + 1;
16. sum = sum + a[j][i + 3] + 1;
17. }
18. }
19. } ...

```

### After Loop Unswitching

```

1. ...
2. int a[N][4];
3. int rand_number = rand();
4. threshold = 2.0 * rand_number;
5. if (threshold < 4) {
6. for (i=0; i<4; i+=4) {
7. for (j=0; j<N; j++) {
8. sum = sum + a[j][i];
9. sum = sum + a[j][i + 1];
10. sum = sum + a[j][i + 2];
11. sum = sum + a[j][i + 3];
12. }
13. }
14. } else {
15. for (i=0; i<4; i+=4) {
16. for (j=0; j<N; j++) {
17. sum = sum + a[j][i] + 1;
18. sum = sum + a[j][i + 1] + 1;
19. sum = sum + a[j][i + 2] + 1;
20. sum = sum + a[j][i + 3] + 1;
21. }
22. }
23. } ...

```

### After Loop Interchange

```

1. ...
2. int a[N][4];
3. int rand_number = rand();
4. threshold = 2.0 * rand_number;
5. if (threshold < 4) {

```

```
6. for (j=0; j<N; j++) {
7.   for (i=0; i<4; i+=4) {
8.     sum = sum + a[j][i];
9.     sum = sum + a[j][i + 1];
10.    sum = sum + a[j][i + 2];
11.    sum = sum + a[j][i + 3];
12.  }
13. }
14. } else {
15.   for (j=0; j<N; j++) {
16.     for (i=0; i<4; i+=4) {
17.       sum = sum + a[j][i] + 1;
18.       sum = sum + a[j][i + 1] + 1;
19.       sum = sum + a[j][i + 2] + 1;
20.       sum = sum + a[j][i + 3] + 1;
21.     }
22.   }
23. } ...
```

## 5. (15 points) Loop Transformations

(a)

**Unsafe** because originally there is **loop carried output dependence** from  $S1[i]$  to  $S3[i - 1]$  and **loop carried anti dependence** from  $S2[i]$  to  $S3[i - 1]$ , however after the loop fusion, the dependences has become **loop carried output dependence** from  $S3[i - 1]$  to  $S1[i]$  and **loop carried true dependence** from  $S3[i-1]$  to  $S2[i]$ .

(b)

**Unsafe** because in the example of loop interchange, the **outermost loop carries a anti dependence** from  $[i-1, j+1]$  to  $[i, j]$ . So here  $i < i'$  and  $j > j'$ , so it is unsafe.

(c) **Safe.**